

Part 5: Structured Support Vector Machines

Sebastian Nowozin and Christoph H. Lampert

Colorado Springs, 25th June 2011

Microsoft

Research



Reminder: learning by regularized risk minimization

For compatibility function $g(x, y; w) := \langle w, \phi(x, y) \rangle$ find w^* that minimizes

$$\mathbb{E}_{(x,y) \sim d(x,y)} \Delta(y, \operatorname{argmax}_y g(x, y; w)).$$

Two major problems:

- ▶ $d(x, y)$ is unknown
- ▶ $\operatorname{argmax}_y g(x, y; w)$ maps into a discrete space
→ $\Delta(y, \operatorname{argmax}_y g(x, y; w))$ is discontinuous, piecewise constant

Task:

$$\min_w \mathbb{E}_{(x,y) \sim d(x,y)} \Delta(y, \operatorname{argmax}_y g(x, y; w)).$$

Problem 1:

- ▶ $d(x, y)$ is unknown

Solution:

- ▶ Replace $\mathbb{E}_{(x,y) \sim d(x,y)}(\cdot)$ with *empirical estimate* $\frac{1}{N} \sum_{(x^n, y^n)}(\cdot)$
- ▶ To avoid overfitting: add a *regularizer*, e.g. $\lambda \|w\|^2$.

New task:

$$\min_w \lambda \|w\|^2 + \frac{1}{N} \sum_{n=1}^N \Delta(y^n, \operatorname{argmax}_y g(x^n, y; w)).$$

Task:

$$\min_w \quad \lambda \|w\|^2 + \frac{1}{N} \sum_{n=1}^N \Delta(y^n, \operatorname{argmax}_y g(x^n, y; w)).$$

Problem:

- ▶ $\Delta(y, \operatorname{argmax}_y g(x, y; w))$ discontinuous w.r.t. w .

Solution:

- ▶ Replace $\Delta(y, y')$ with *well behaved* $\ell(x, y, w)$
- ▶ Typically: ℓ *upper bound* to Δ , *continuous* and *convex* w.r.t. w .

New task:

$$\min_w \quad \lambda \|w\|^2 + \frac{1}{N} \sum_{n=1}^N \ell(x^n, y^n, w)$$

Regularized Risk Minimization

$$\min_w \quad \lambda \|w\|^2 \quad + \quad \frac{1}{N} \sum_{n=1}^N \ell(x^n, y^n, w)$$

Regularization + Loss on training data

Hinge loss: maximum margin training

$$\ell(x^n, y^n, w) := \max_{y \in \mathcal{Y}} \left[\Delta(y^n, y) + \langle w, \phi(x^n, y) \rangle - \langle w, \phi(x^n, y^n) \rangle \right]$$

Alternative:

Logistic loss: probabilistic training

$$\ell(x^n, y^n, w) := \log \sum_{y \in \mathcal{Y}} \exp \left(\langle w, \phi(x^n, y) \rangle - \langle w, \phi(x^n, y^n) \rangle \right)$$

Structured Output Support Vector Machine

$$\min_w \frac{1}{2} \|w\|^2 + \frac{C}{N} \sum_{n=1}^N \left[\max_{y \in \mathcal{Y}} \Delta(y^n, y) + \langle w, \phi(x^n, y) \rangle - \langle w, \phi(x^n, y^n) \rangle \right]$$

Conditional Random Field

$$\min_w \frac{\|w\|^2}{2\sigma^2} + \sum_{n=1}^N \left[\log \sum_{y \in \mathcal{Y}} \exp(\langle w, \phi(x^n, y) \rangle - \langle w, \phi(x^n, y^n) \rangle) \right]$$

CRFs and SSVMs have more in common than usually assumed.

- ▶ both do regularized risk minimization
- ▶ $\log \sum_y \exp(\cdot)$ can be interpreted as a *soft-max*

Solving the Training Optimization Problem Numerically

Structured Output Support Vector Machine:

$$\min_w \frac{1}{2} \|w\|^2 + \frac{C}{N} \sum_{n=1}^N \left[\max_{y \in \mathcal{Y}} \Delta(y^n, y) + \langle w, \phi(x^n, y) \rangle - \langle w, \phi(x^n, y^n) \rangle \right]$$

Unconstrained optimization, convex, non-differentiable objective.

Structured Output SVM (equivalent formulation):

$$\min_{w, \xi} \quad \frac{1}{2} \|w\|^2 + \frac{C}{N} \sum_{n=1}^N \xi^n$$

subject to, for $n = 1, \dots, N$,

$$\max_{y \in \mathcal{Y}} \left[\Delta(y^n, y) + \langle w, \phi(x^n, y) \rangle - \langle w, \phi(x^n, y^n) \rangle \right] \leq \xi^n$$

N non-linear constraints, convex, differentiable objective.

Structured Output SVM (also equivalent formulation):

$$\min_{w, \xi} \quad \frac{1}{2} \|w\|^2 + \frac{C}{N} \sum_{n=1}^N \xi^n$$

subject to, for $n = 1, \dots, N$,

$$\Delta(y^n, y) + \langle w, \phi(x^n, y) \rangle - \langle w, \phi(x^n, y^n) \rangle \leq \xi^n, \quad \text{for all } y \in \mathcal{Y}$$

$N|\mathcal{Y}|$ linear constraints, convex, differentiable objective.

Summary – S-SVM Learning

Given:

- ▶ training set $\{(x^1, y^1), \dots, (x^n, y^n)\} \subset \mathcal{X} \times \mathcal{Y}$
- ▶ loss function $\Delta : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$.

Task: learn parameter w for $f(x) := \operatorname{argmax}_y \langle w, \phi(x, y) \rangle$ that minimizes expected loss on future data.

S-SVM solution derived by *maximum margin* framework:

- ▶ enforce **correct output** to be better than **others** by a **margin**:

$$\langle w, \phi(x^n, y^n) \rangle \geq \Delta(y^n, y) + \langle w, \phi(x^n, y) \rangle \quad \text{for all } y \in \mathcal{Y}.$$

- ▶ convex optimization problem, but non-differentiable
- ▶ many equivalent formulations \rightarrow different training algorithms
- ▶ training needs repeated argmax prediction, no probabilistic inference

Solving the Training Optimization Problem Numerically

We can solve an S-SVM like a linear SVM:

One of the equivalent formulations was:

$$\min_{w \in \mathbb{R}^D, \xi \in \mathbb{R}_+^n} \|w\|^2 + \frac{C}{N} \sum_{n=1}^N \xi^n$$

subject to, for $i = 1, \dots, n$,

$$\langle w, \phi(x^n, y^n) \rangle - \langle w, \phi(x^n, y) \rangle \geq \Delta(y^n, y) - \xi^n, \quad \text{for all } y \in \mathcal{Y}'.$$

Introduce feature vectors $\delta\phi(x^n, y^n, y) := \phi(x^n, y^n) - \phi(x^n, y)$.

Solve

$$\min_{w \in \mathbb{R}^D, \xi \in \mathbb{R}_+^n} \|w\|^2 + \frac{C}{N} \sum_{n=1}^N \xi^n$$

subject to, for $i = 1, \dots, n$, for all $y \in \mathcal{Y}$,

$$\langle w, \delta\phi(x^n, y^n, y) \rangle \geq \Delta(y^n, y) - \xi^n.$$

This has the same structure as an ordinary SVM!

- ▶ quadratic objective ☺
- ▶ linear constraints ☺

Solve

$$\min_{w \in \mathbb{R}^D, \xi \in \mathbb{R}_+^n} \|w\|^2 + \frac{C}{N} \sum_{n=1}^N \xi^n$$

subject to, for $i = 1, \dots, n$, for all $y \in \mathcal{Y}$,

$$\langle w, \delta\phi(x^n, y^n, y) \rangle \geq \Delta(y^n, y) - \xi^n.$$

This has the same structure as an ordinary SVM!

- ▶ quadratic objective ☺
- ▶ linear constraints ☺

Question: Can't we use a ordinary SVM/QP solver?

Solve

$$\min_{w \in \mathbb{R}^D, \xi \in \mathbb{R}_+^n} \|w\|^2 + \frac{C}{N} \sum_{n=1}^N \xi^n$$

subject to, for $i = 1, \dots, n$, for all $y \in \mathcal{Y}$,

$$\langle w, \delta\phi(x^n, y^n, y) \rangle \geq \Delta(y^n, y) - \xi^n.$$

This has the same structure as an ordinary SVM!

- ▶ quadratic objective ☺
- ▶ linear constraints ☺

Question: Can't we use a ordinary SVM/QP solver?

Answer: Almost! We could, if there weren't $N|\mathcal{Y}|$ constraints.

- ▶ E.g. 100 binary 16×16 images: 10^{79} constraints

Solution: working set training

- ▶ It's enough if we enforce the **active constraints**.
The others will be fulfilled automatically.
- ▶ We don't know which ones are active for the optimal solution.
- ▶ But it's likely to be only a small number \leftarrow can of course be formalized.

Keep a set of potentially active constraints and update it iteratively:

Solution: working set training

- ▶ It's enough if we enforce the **active constraints**.
The others will be fulfilled automatically.
- ▶ We don't know which ones are active for the optimal solution.
- ▶ But it's likely to be only a small number ← can of course be formalized.

Keep a set of potentially active constraints and update it iteratively:

Working Set Training

- ▶ Start with working set $S = \emptyset$ (no constraints)
- ▶ Repeat until convergence:
 - ▶ Solve S-SVM training problem with constraints from S
 - ▶ Check, if solution violates any of the *full* constraint set
 - ▶ if no: we found the optimal solution, *terminate*.
 - ▶ if yes: add most violated constraints to S , *iterate*.

Solution: working set training

- ▶ It's enough if we enforce the **active constraints**.
The others will be fulfilled automatically.
- ▶ We don't know which ones are active for the optimal solution.
- ▶ But it's likely to be only a small number \leftarrow can of course be formalized.

Keep a set of potentially active constraints and update it iteratively:

Working Set Training

- ▶ Start with working set $S = \emptyset$ (no constraints)
- ▶ Repeat until convergence:
 - ▶ Solve S-SVM training problem with constraints from S
 - ▶ Check, if solution violates any of the *full* constraint set
 - ▶ if no: we found the optimal solution, *terminate*.
 - ▶ if yes: add most violated constraints to S , *iterate*.

Good *practical performance* and *theoretic guarantees*:

- ▶ polynomial time convergence ϵ -close to the global optimum

Working Set S-SVM Training

input training pairs $\{(x^1, y^1), \dots, (x^n, y^n)\} \subset \mathcal{X} \times \mathcal{Y}$,

input feature map $\phi(x, y)$, loss function $\Delta(y, y')$, regularizer C

```
1:  $S \leftarrow \emptyset$ 
2: repeat
3:    $(w, \xi) \leftarrow \text{solution to QP only with constraints from } S$ 
4:   for  $i=1, \dots, n$  do
5:      $\hat{y} \leftarrow \operatorname{argmax}_{y \in \mathcal{Y}} \Delta(y^n, y) + \langle w, \phi(x^n, y) \rangle$ 
6:     if  $\hat{y} \neq y^n$  then
7:        $S \leftarrow S \cup \{(x^n, \hat{y})\}$ 
8:     end if
9:   end for
10: until  $S$  doesn't change anymore.
```

output prediction function $f(x) = \operatorname{argmax}_{y \in \mathcal{Y}} \langle w, \phi(x, y) \rangle$.

Observation: each update of w needs 1 argmax -prediction per example.
(but we solve globally for next w , not by local steps)